A very brief Introduction to Predictive Models for Data Science

Introduction to Data Science

Slides by Mark Carman

Motivating example

(Example from Duda & Hart, Pattern Classification & Scene Analysis, 1973)

On a fishing boat, a conveyor belt loads fish and a worker separates them into their different types

Motivating example

(Example from Duda & Hart, Pattern Classification & Scene Analysis, 1973)

On a fishing boat, a conveyor belt loads fish and a worker separates them into their different types

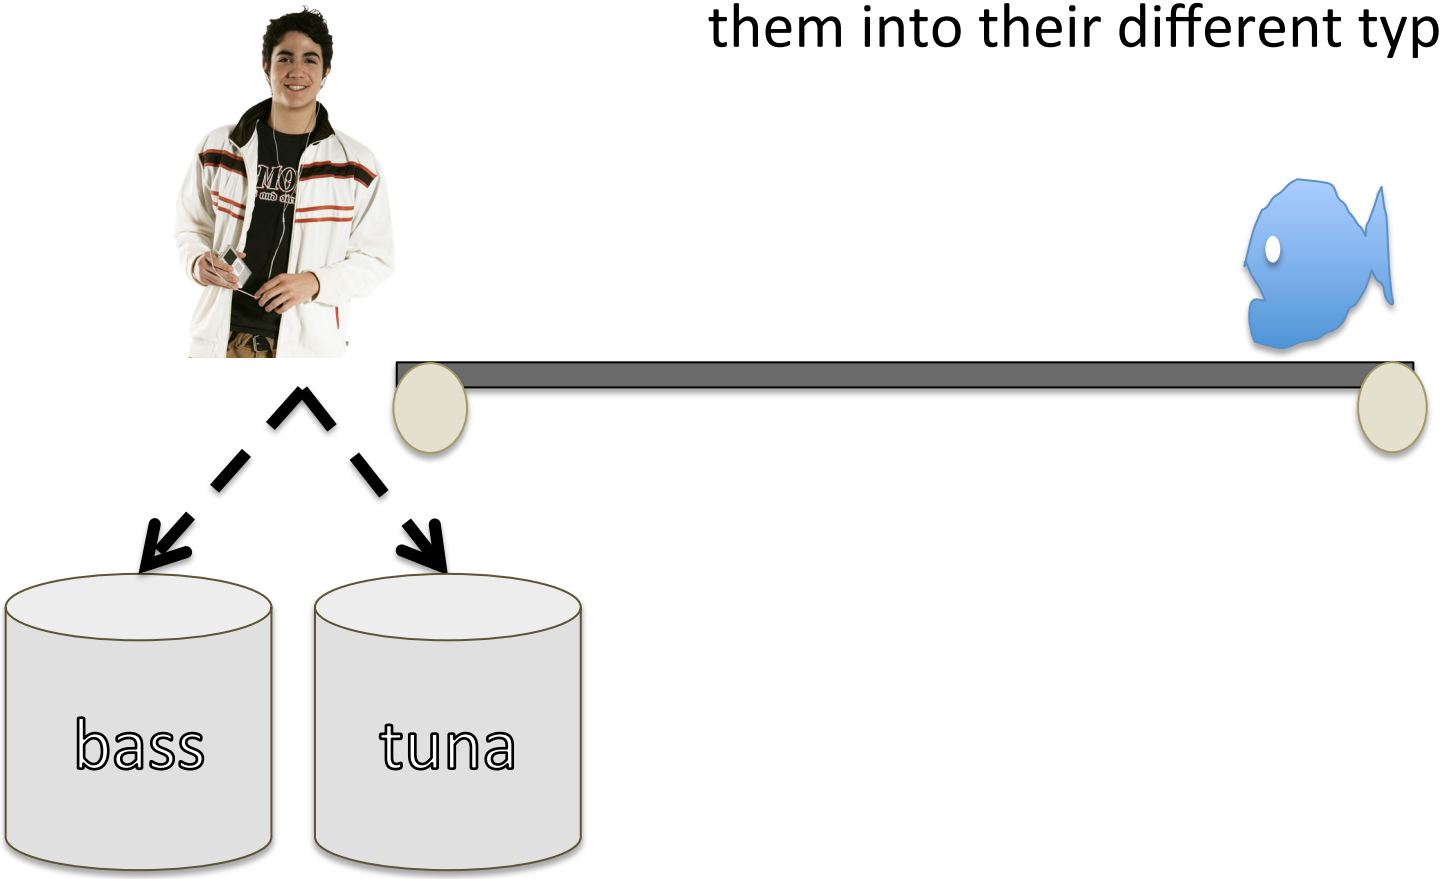

Motivating example

On a fishing boat, a conveyor belt loads fish and a worker separates them into their different types

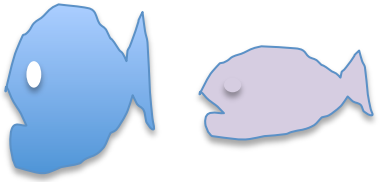

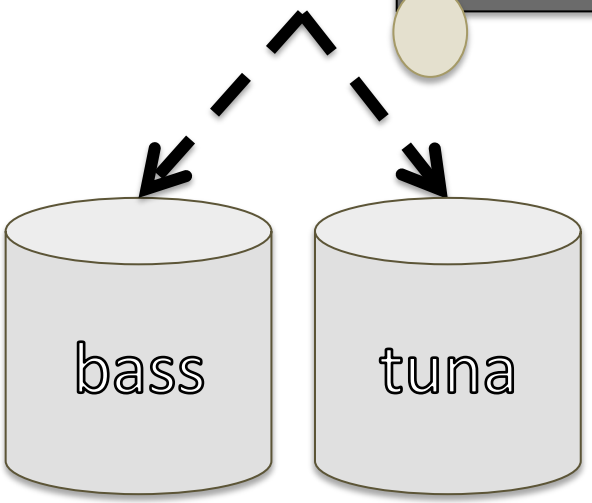

Motivating example

On a fishing boat, a conveyor belt loads fish and a worker separates them into their different types

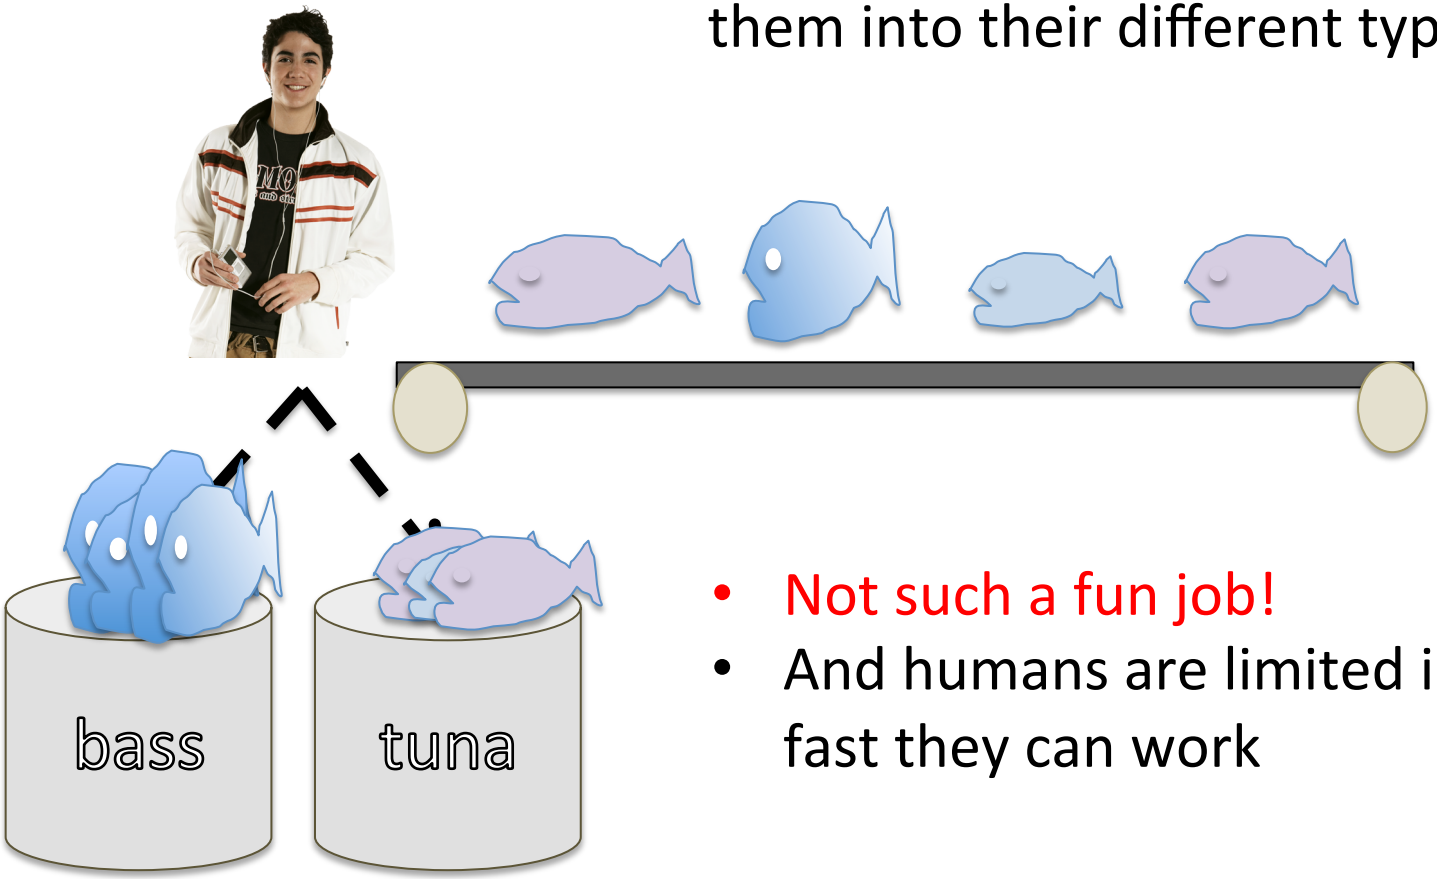

- Not such a fun job!
- And humans are limited in how fast they can work

Motivating example

Question: Can we build a system to do the task automatically?

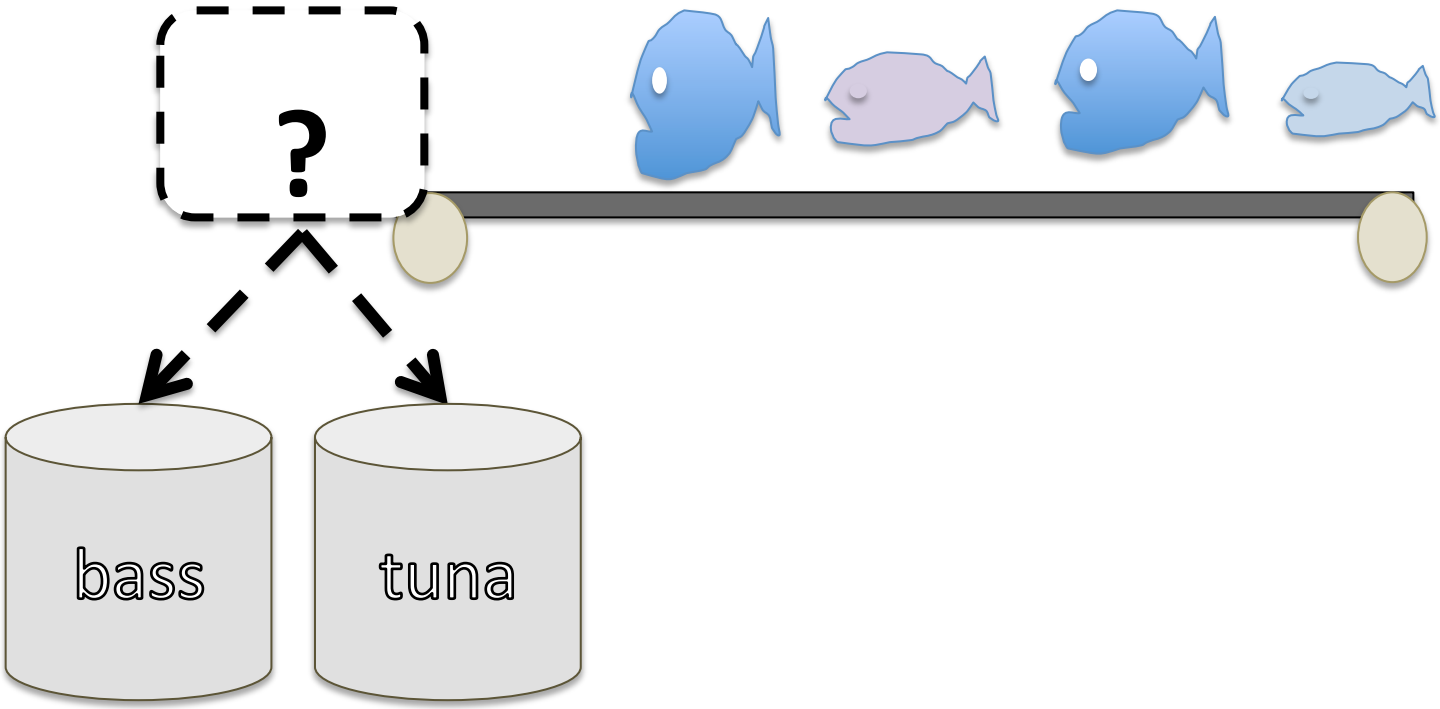

Motivating example

1. Install camera to capture image of fish
2. Compute **features** of fish from image:
   - X_1: length
   - X_2: width
   - X_3: colour

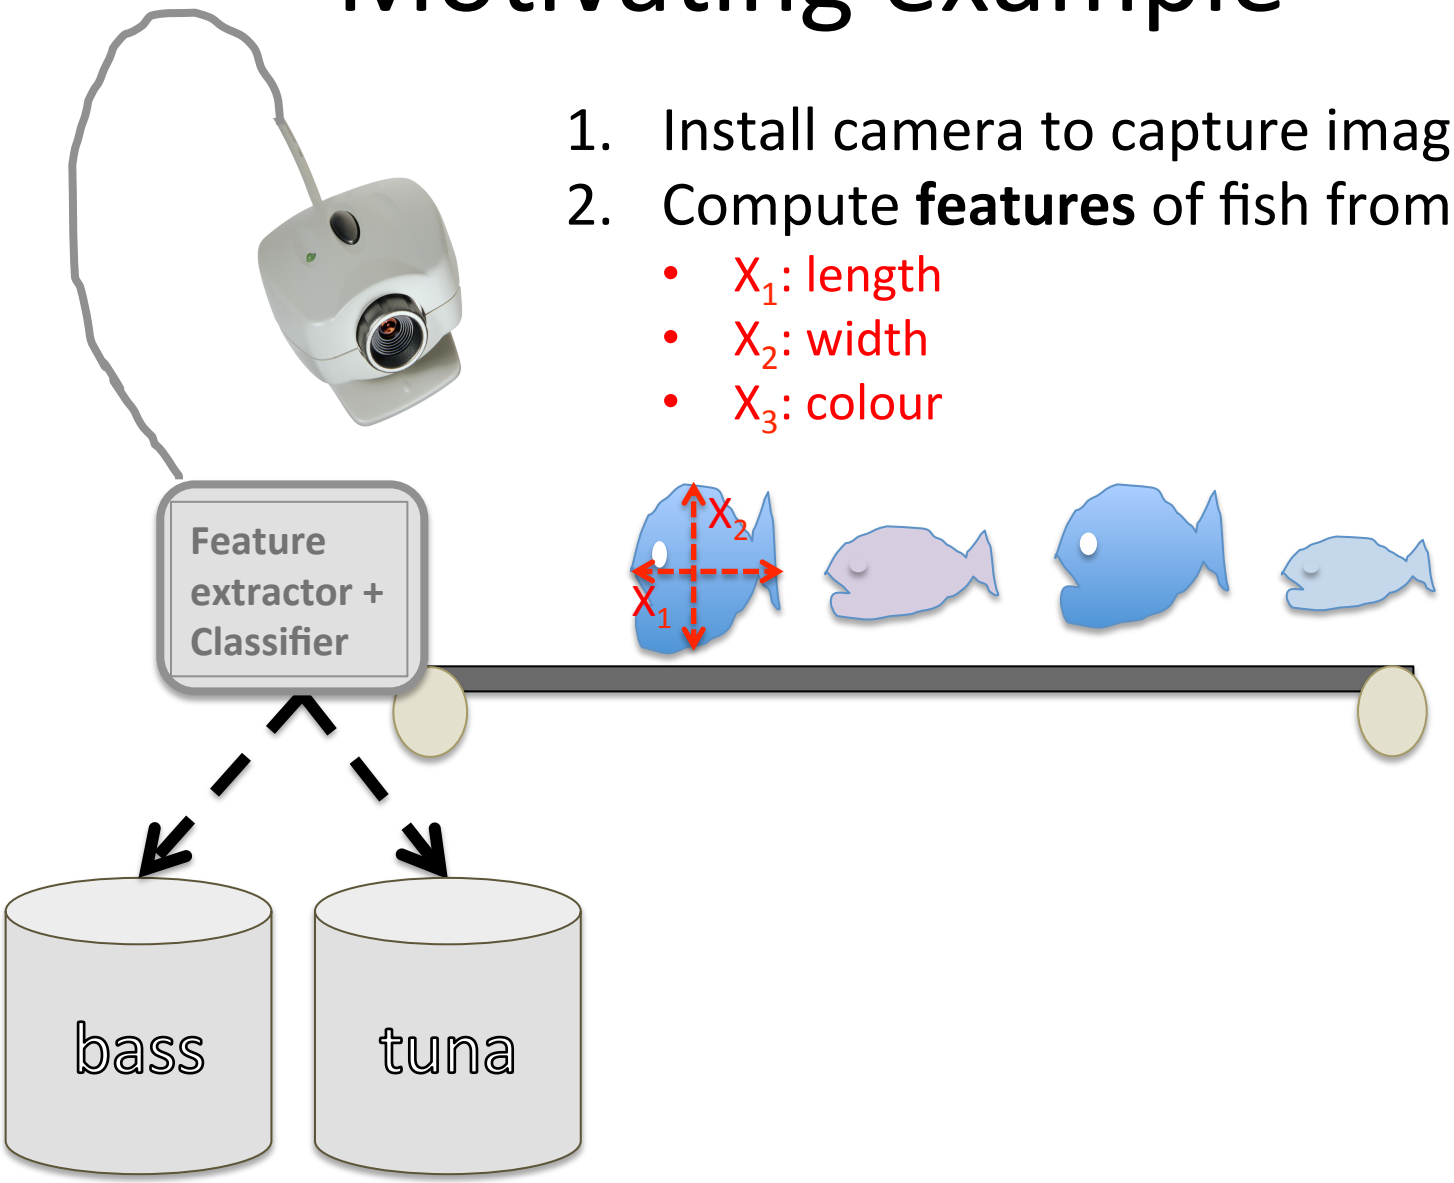

Motivating example

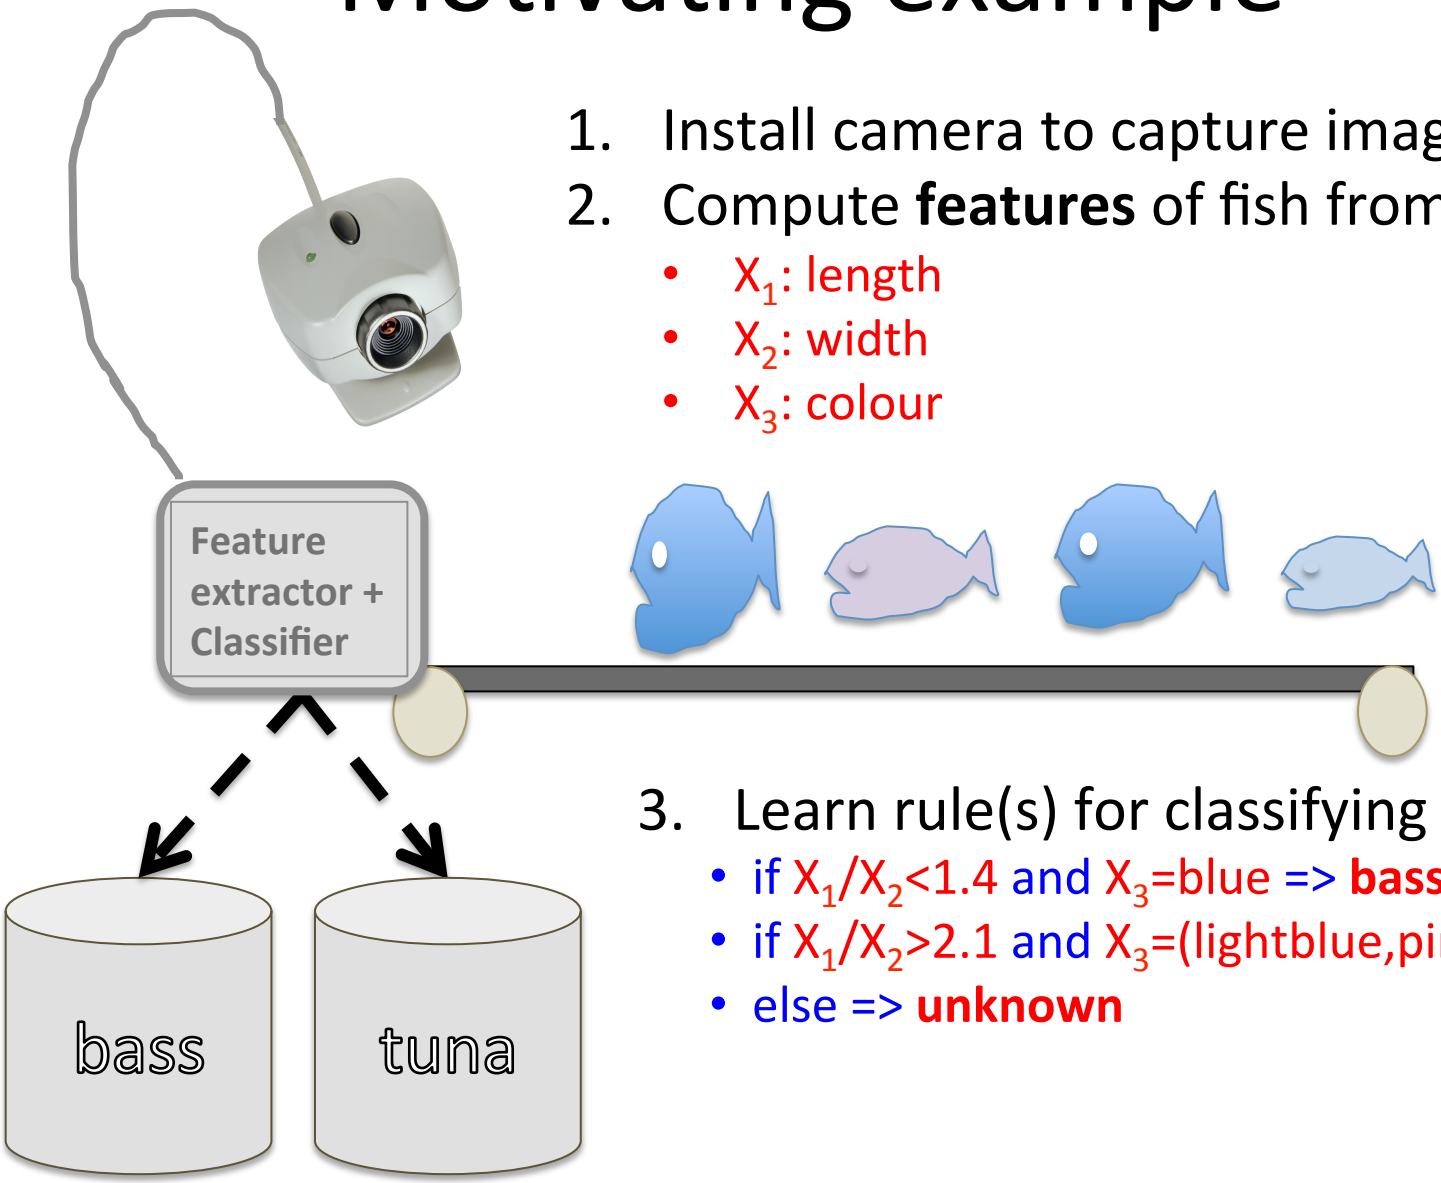

1. Install camera to capture image of fish
2. Compute **features** of fish from image:
   - X_1: length
   - X_2: width
   - X_3: colour

3. Learn rule(s) for classifying fish
   - if $X_1/X_2<1.4$ and X_3=blue => **bass**
   - if $X_1/X_2>2.1$ and X_3=(lightblue,pink) => **tuna**
   - else => **unknown**

Predictive Models

A predictive model is any model that makes a prediction

- usually based on a set of features describing an object.

The prediction could be:

- a binary outcome (spam, not-spam)
- categorical (bass, tuna, other)
- a real value (the age of the fish)
- a vector of real values (probability of bass, tuna, etc.)
- Etc.

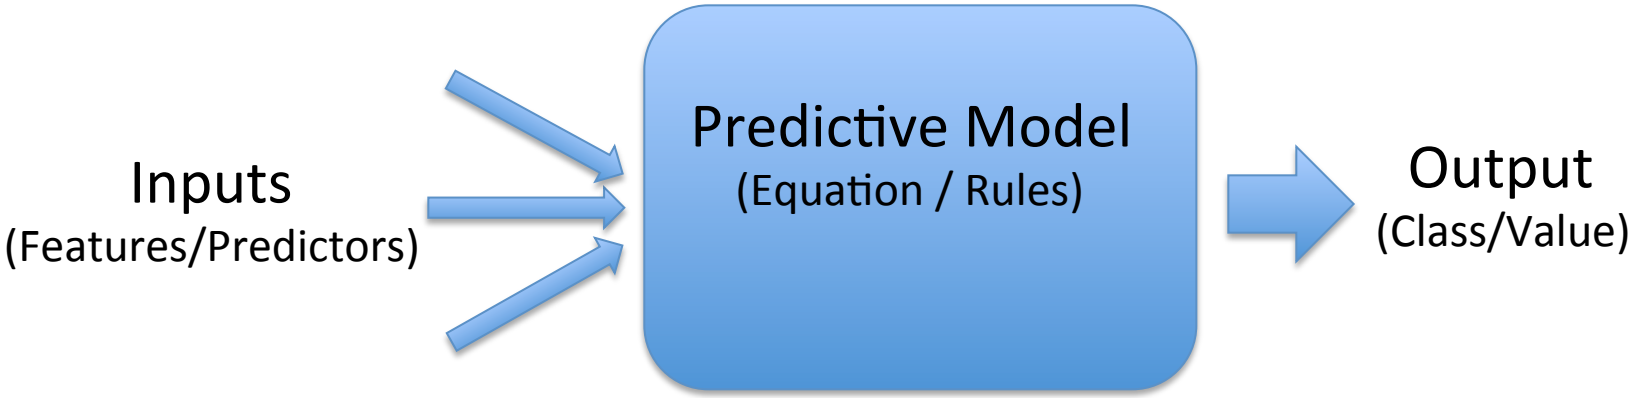

Predictive Models

- If the predicted value is binary/categorical we usually refer to the model as a classifier
- If it predicts real values we refer to it as regression
- Although there are many other types of model (e.g. ranking, translation, etc.)

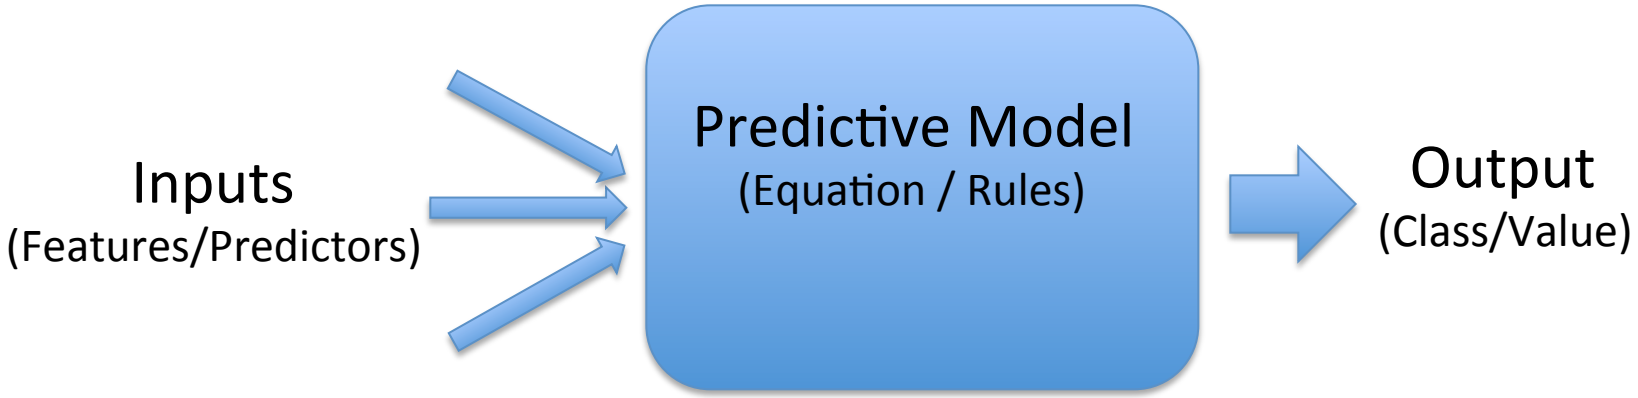

Predictive Models

The predictive model uses equations/rules to map the input features to output values

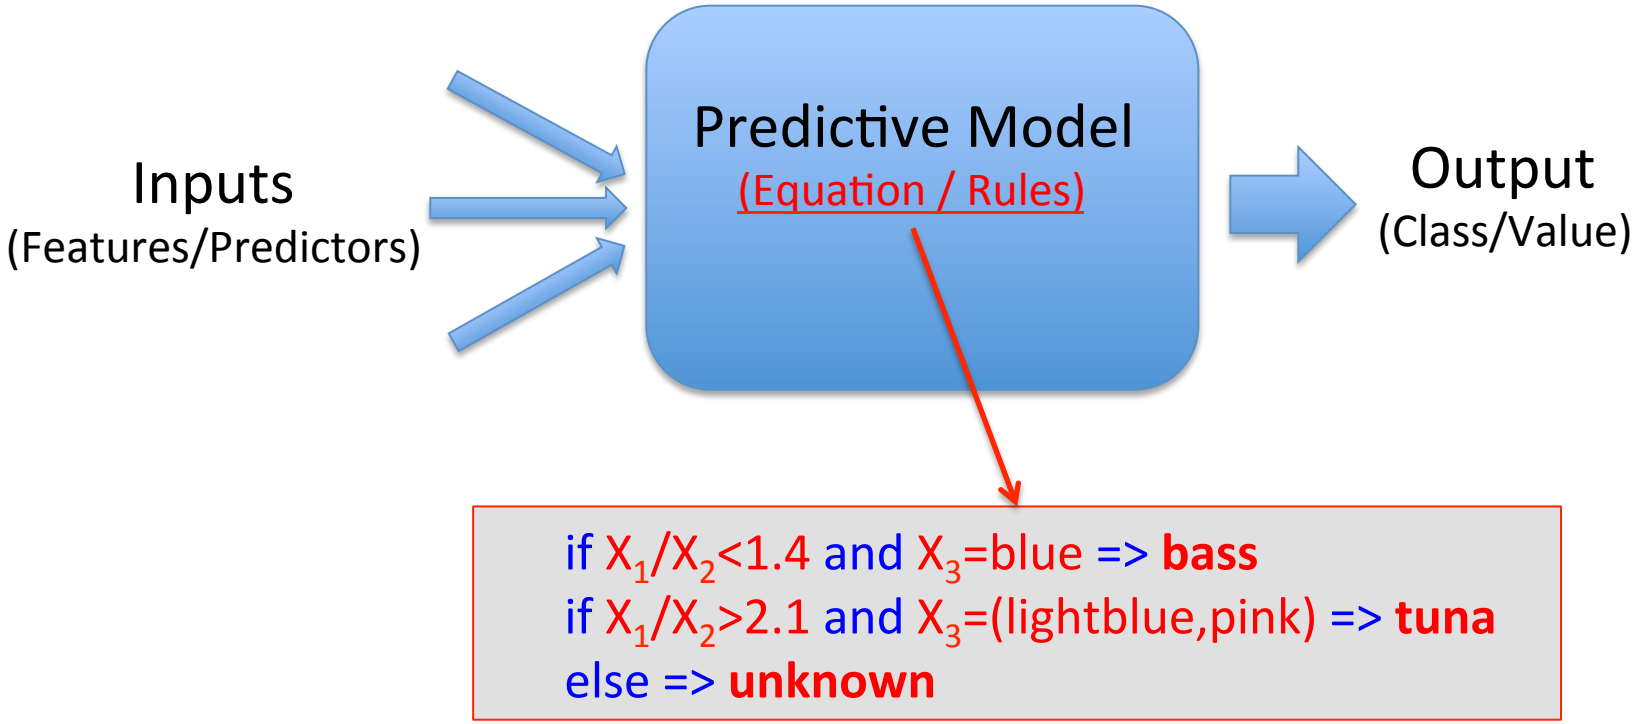

Predictive Model

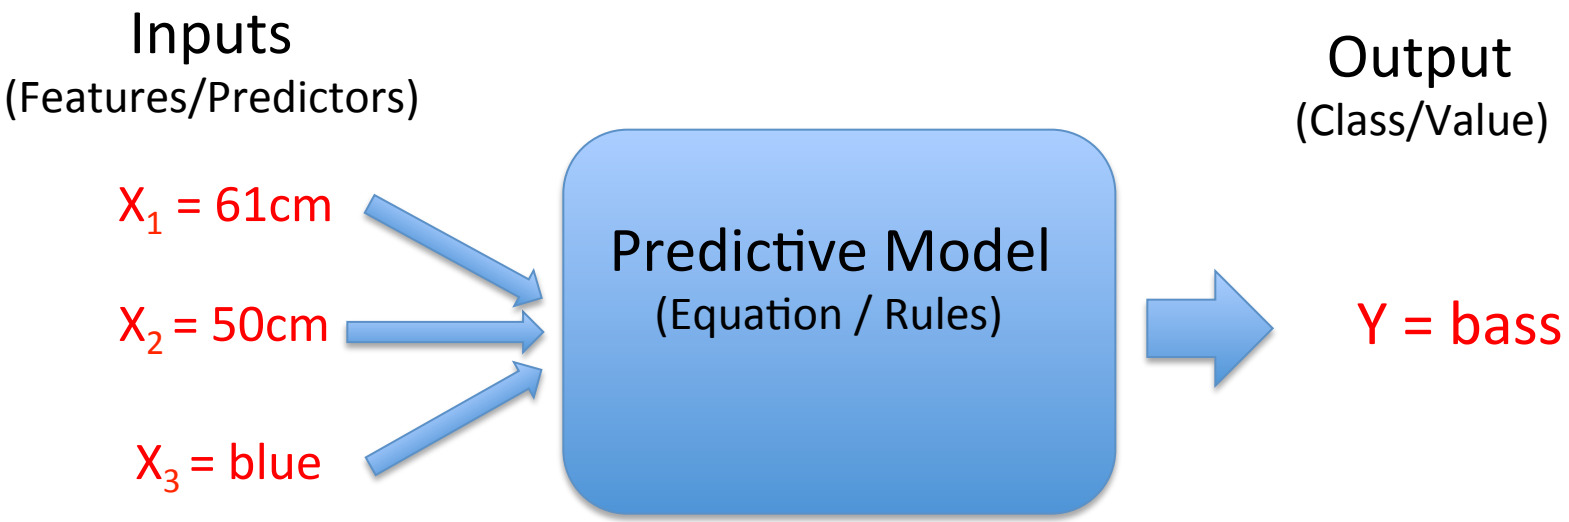

Predictive Model

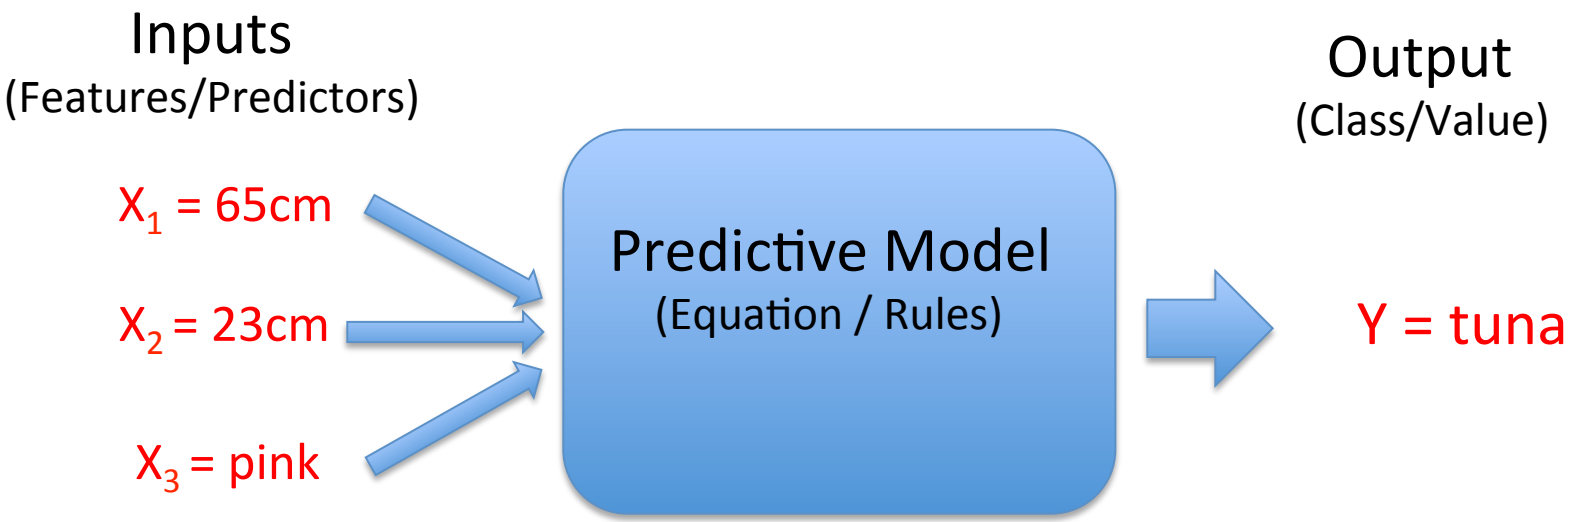

Models are learnt from Examples

Instance	X1 = length	X2 = width	X3 = colour	Y = class
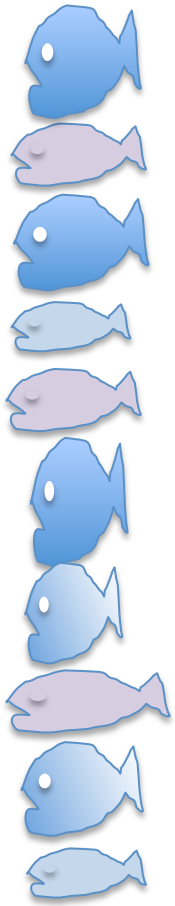	55	51	blue	**bass**
	65	23	pink	**tuna**
	67	54	blue	**bass**
	54	20	light-blue	**tuna**
	62	26	pink	**tuna**
	44	62	blue	**bass**
	47	55	light-blue	**bass**
	73	31	pink	**tuna**
	54	48	light-blue	**bass**
	57	23	light-blue	**tuna**

Training a Model

Predictive models are learnt from training data and then applied to make predictions on new instances

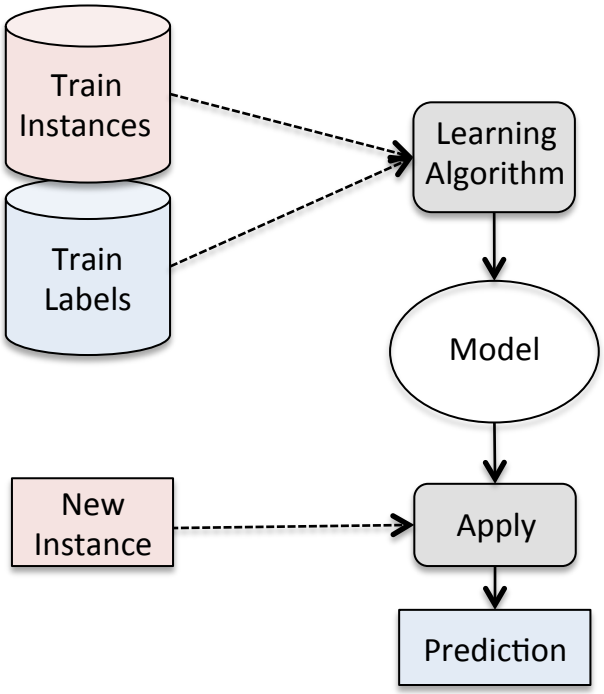

How are models learnt?

- Each training instance (fish in our case) is just a point in some feature space
- Here the colour denotes the class
  - (blue = bass, green = tuna, red = unknown)

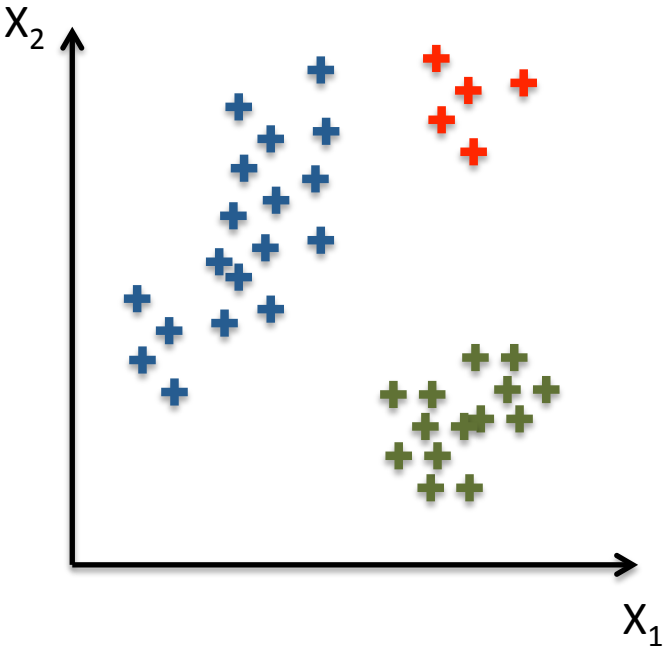

How are models learnt?

Many (classification) learning algorithms work by dividing the feature space into regions of the same type

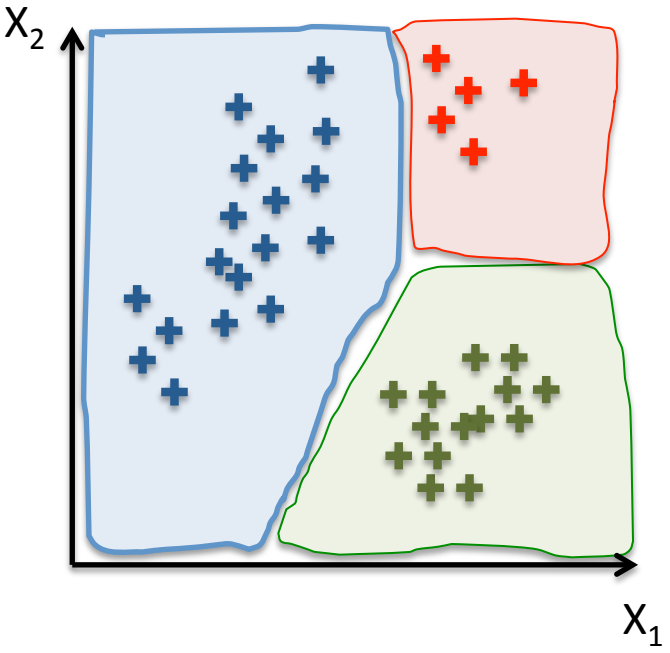

In practice

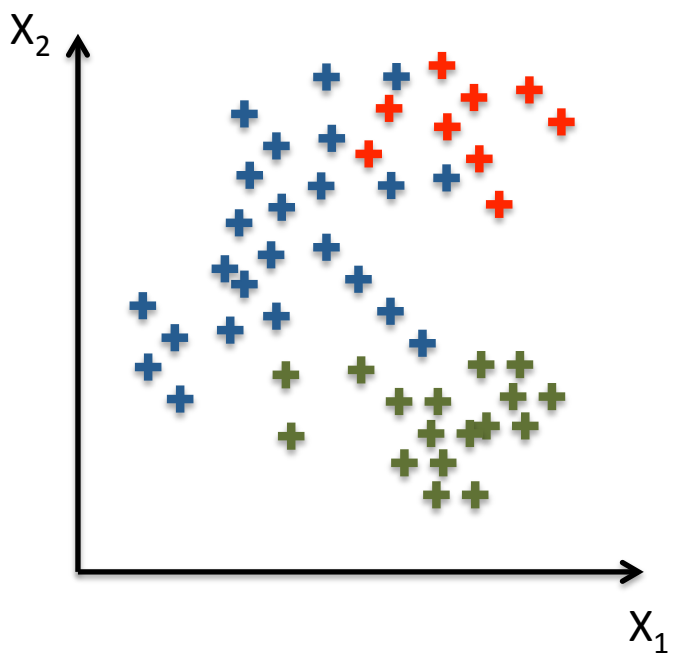

- In practice, the data is usually overlapping
- making it hard to separate the classes

In practice

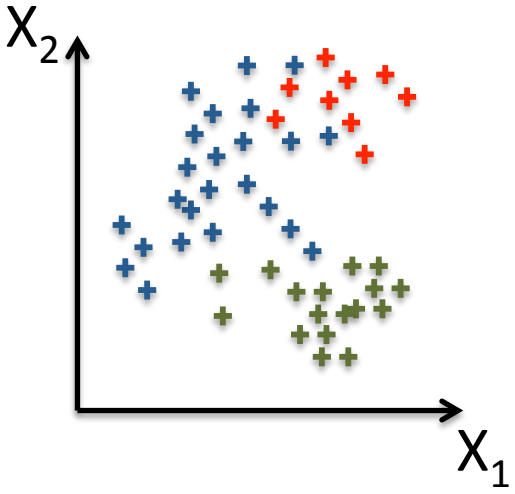

- and we have many feature dimensions
- with some features more useful than others

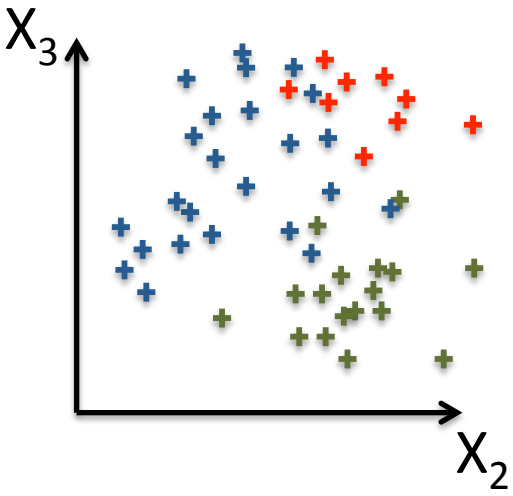

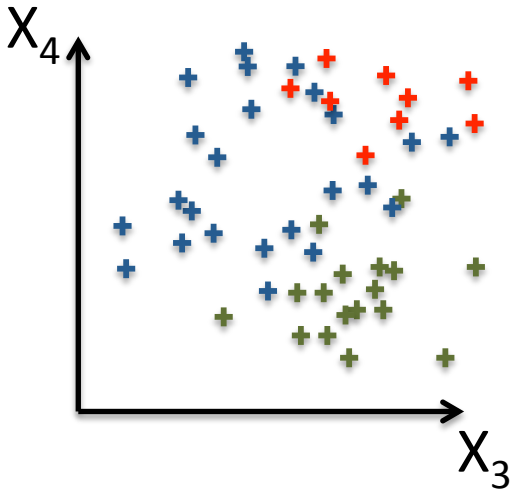

Different Models

There are many different types of models that we can train to classify objects

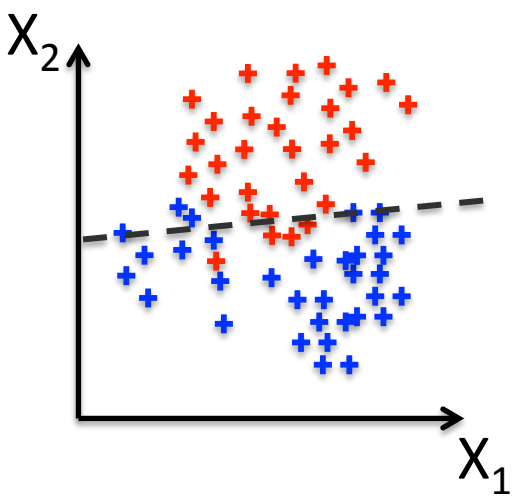

Linear classifiers

e.g. Logistic Regression, Linear SVMs

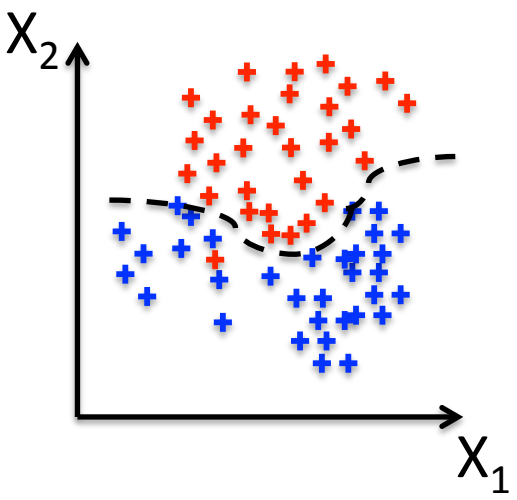

Non-linear Classifiers

e.g. Neural Nets, SVM with RBF kernel

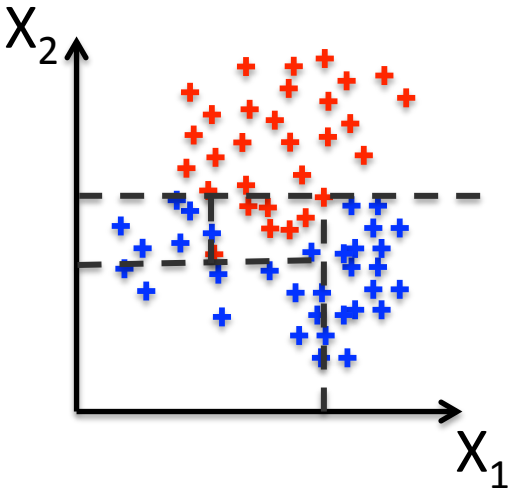

Decision Tree Learners

e.g. Random forests

Testing models

We evaluate predictive models based on how well they predict the labels for test instances (not used in training)

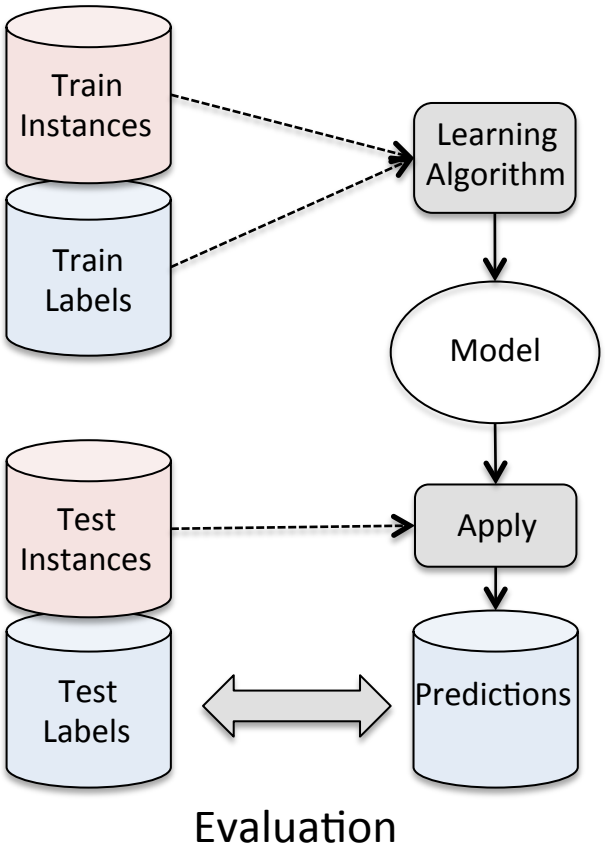

Performance of predictive models

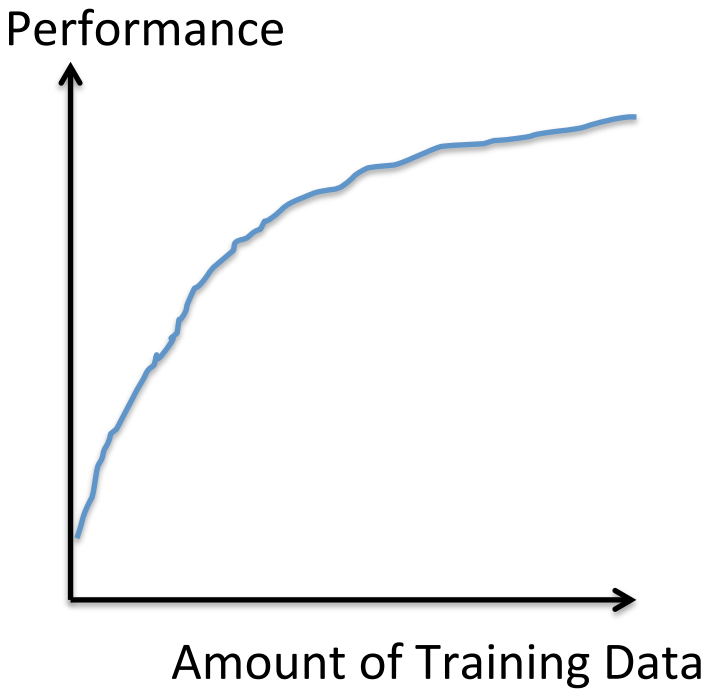

Generally:

- The more training data the better the test performance
- And (providing there is sufficient training data) the more features the better performance

End of Introduction!

- We'll talk more about predictive models in the coming weeks, especially in module 5.